

2203a51706-ml-assign

March 5, 2024

1 2203A51706 ML ASSIGNMENT

```
[1]: import pandas as pd
import seaborn as sns
import os
import numpy as np
import matplotlib.pyplot as plt
```

```
[ ]: housing_df = pd.read_csv('/content/housing.csv')

# Use .info() to show the features (i.e. columns) in your dataset along with a
↳ count and datatype
housing_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
4   total_bedrooms         20433 non-null  float64
5   population             20640 non-null  float64
6   households              20640 non-null  float64
7   median_income          20640 non-null  float64
8   median_house_value     20640 non-null  float64
9   ocean_proximity        20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

```
[3]: housing_df.shape
```

```
[3]: (20640, 10)
```

```
[4]: housing_df.head()
```

```
[4]: longitude latitude housing_median_age total_rooms total_bedrooms \
0 -122.23 37.88 41.0 880.0 129.0
1 -122.22 37.86 21.0 7099.0 1106.0
2 -122.24 37.85 52.0 1467.0 190.0
3 -122.25 37.85 52.0 1274.0 235.0
4 -122.25 37.85 52.0 1627.0 280.0

population households median_income median_house_value ocean_proximity
0 322.0 126.0 8.3252 452600.0 NEAR BAY
1 2401.0 1138.0 8.3014 358500.0 NEAR BAY
2 496.0 177.0 7.2574 352100.0 NEAR BAY
3 558.0 219.0 5.6431 341300.0 NEAR BAY
4 565.0 259.0 3.8462 342200.0 NEAR BAY
```

```
[5]: housing_df.tail()
```

```
[5]: longitude latitude housing_median_age total_rooms total_bedrooms \
20635 -121.09 39.48 25.0 1665.0 374.0
20636 -121.21 39.49 18.0 697.0 150.0
20637 -121.22 39.43 17.0 2254.0 485.0
20638 -121.32 39.43 18.0 1860.0 409.0
20639 -121.24 39.37 16.0 2785.0 616.0

population households median_income median_house_value \
20635 845.0 330.0 1.5603 78100.0
20636 356.0 114.0 2.5568 77100.0
20637 1007.0 433.0 1.7000 92300.0
20638 741.0 349.0 1.8672 84700.0
20639 1387.0 530.0 2.3886 89400.0

ocean_proximity
20635 INLAND
20636 INLAND
20637 INLAND
20638 INLAND
20639 INLAND
```

```
[6]: housing_df.describe()
```

```
[6]: longitude latitude housing_median_age total_rooms \
count 20640.000000 20640.000000 20640.000000 20640.000000
mean -119.569704 35.631861 28.639486 2635.763081
std 2.003532 2.135952 12.585558 2181.615252
min -124.350000 32.540000 1.000000 2.000000
25% -121.800000 33.930000 18.000000 1447.750000
50% -118.490000 34.260000 29.000000 2127.000000
75% -118.010000 37.710000 37.000000 3148.000000
```

```
max      -114.310000      41.950000      52.000000  39320.000000
```

	total_bedrooms	population	households	median_income \
count	20433.000000	20640.000000	20640.000000	20640.000000
mean	537.870553	1425.476744	499.539680	3.870671
std	421.385070	1132.462122	382.329753	1.899822
min	1.000000	3.000000	1.000000	0.499900
25%	296.000000	787.000000	280.000000	2.563400
50%	435.000000	1166.000000	409.000000	3.534800
75%	647.000000	1725.000000	605.000000	4.743250
max	6445.000000	35682.000000	6082.000000	15.000100

	median_house_value
count	20640.000000
mean	206855.816909
std	115395.615874
min	14999.000000
25%	119600.000000
50%	179700.000000
75%	264725.000000
max	500001.000000

```
[7]: housing_df.isnull().sum()
```

```
[7]: longitude      0
latitude      0
housing_median_age  0
total_rooms    0
total_bedrooms 207
population     0
households     0
median_income  0
median_house_value  0
ocean_proximity  0
dtype: int64
```

```
[8]: # Calculate the % of missing data
housing_df['total_bedrooms'].isnull().sum()/housing_df.shape[0] * 100
```

```
[8]: 1.002906976744186
```

```
[9]: from sklearn.impute import KNNImputer

# create a temporary copy of the dataset
housing_df_temp = housing_df.copy()
```

```

# retrieve columns with numerical data; will exclude the ocean_proximity column
↳since the datatype is object; other columns are float64
columns_list = [col for col in housing_df_temp.columns if housing_df_temp[col].
↳dtype != 'object']

# extract columns that contain at least one missing value
new_column_list = [col for col in housing_df_temp.loc[:, housing_df_temp.
↳isnull().any()]]

# update temp dataframe with numeric columns that have empty values
housing_df_temp = housing_df_temp[new_column_list]

```

```

[10]: # initialize KNNImputer to impute missing data using machine learning
knn = KNNImputer(n_neighbors = 3)

# fit function trains the model
knn.fit(housing_df_temp)

# transform the data using the model
# applies the transformation model (ie knn) to data
array_Values = knn.transform(housing_df_temp)

# convert the array values to a dataframe with the appropriate column names
housing_df_temp = pd.DataFrame(array_Values, columns = new_column_list)

```

```

[11]: # confirm there are no columns with missing values
housing_df_temp.isnull().sum()

```

```

[11]: total_bedrooms    0
dtype: int64

```

```

[12]: # overlay the imputed column over the old column with missing values

# loop through the list of columns and overlay each one
for column_name in new_column_list:
    housing_df[column_name] = housing_df_temp.
↳replace(housing_df[column_name],housing_df[column_name])

# confirm columns no longer contain null data
housing_df.isnull().sum()

```

```

[12]: longitude          0
latitude                0
housing_median_age      0
total_rooms              0
total_bedrooms          0
population              0

```


```
households      0
median_income    0
median_house_value  0
ocean_proximity  0
dtype: int64
```

```
[13]: # Plot the distribution of the target variable (median_house_value) using a
      ↪ histogram

      # bins->amount of columns
      plt.hist(housing_df['median_house_value'], bins=80)
      plt.xlabel("House Values")

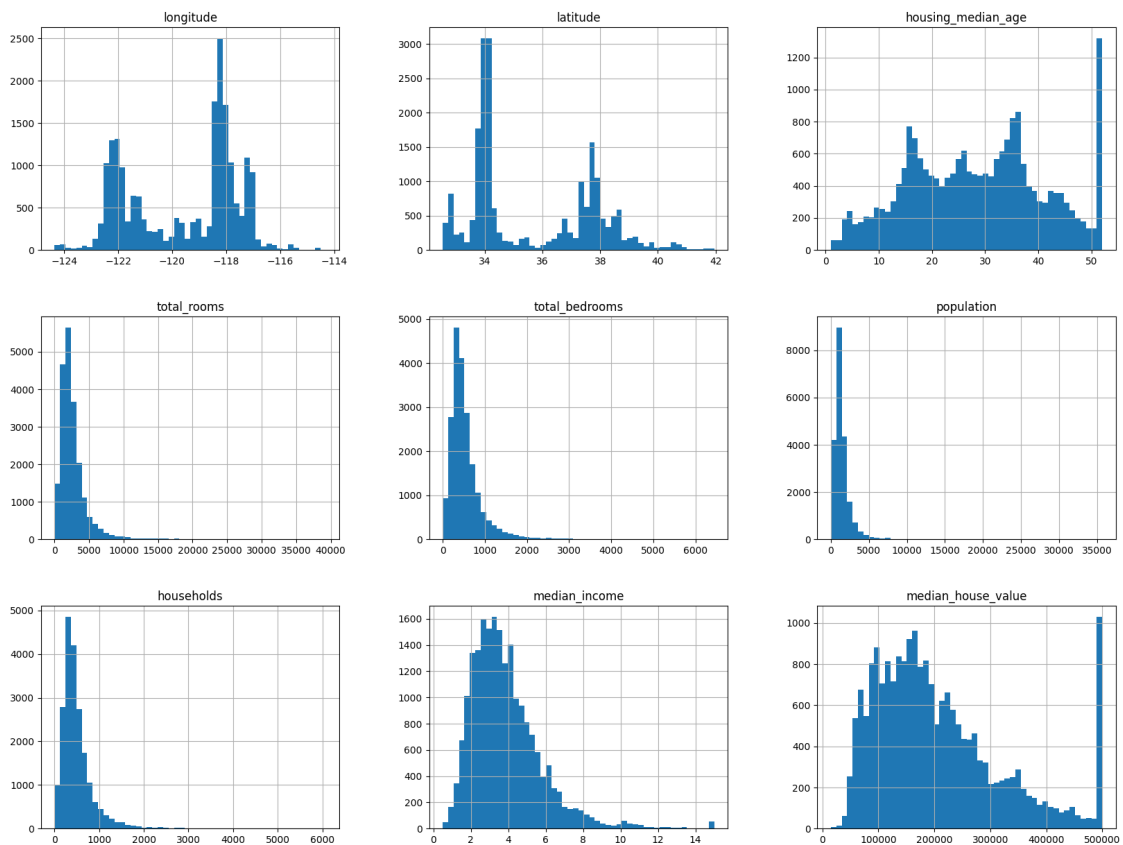
      # We can see from the plot that the values of Median House Value are
      ↪ distributed normally with few outliers.
      # Most of the house are around 100,000-200,000 range
```

```
[13]: Text(0.5, 0, 'House Values')
```



```
[14]: # let's do histograms for the all the features to understand the data_
      ↪distributions
      # using housing_df as to not plot the encoded values for OCEAN_PROXIMITY
      housing_df.hist(bins=50, figsize=(20,15))
```

```
[14]: array([[<Axes: title={'center': 'longitude'}>,
      <Axes: title={'center': 'latitude'}>,
      <Axes: title={'center': 'housing_median_age'}>],
      [<Axes: title={'center': 'total_rooms'}>,
      <Axes: title={'center': 'total_bedrooms'}>,
      <Axes: title={'center': 'population'}>],
      [<Axes: title={'center': 'households'}>,
      <Axes: title={'center': 'median_income'}>,
      <Axes: title={'center': 'median_house_value'}>]], dtype=object)
```



```
[15]: # Plot a graphical correlation matrix for each pair of columns in the dataframe
      corr = housing_df.corr() # data frame correlation function
      print(corr)
```

	longitude	latitude	housing_median_age	total_rooms \
longitude	1.000000	-0.924664	-0.108197	0.044568

latitude	-0.924664	1.000000	0.011173	-0.036100
housing_median_age	-0.108197	0.011173	1.000000	-0.361262
total_rooms	0.044568	-0.036100	-0.361262	1.000000
total_bedrooms	0.069260	-0.066658	-0.318998	0.927253
population	0.099773	-0.108785	-0.296244	0.857126
households	0.055310	-0.071035	-0.302916	0.918484
median_income	-0.015176	-0.079809	-0.119034	0.198050
median_house_value	-0.045967	-0.144160	0.105623	0.134153

	total_bedrooms	population	households	median_income	\
longitude	0.069260	0.099773	0.055310	-0.015176	
latitude	-0.066658	-0.108785	-0.071035	-0.079809	
housing_median_age	-0.318998	-0.296244	-0.302916	-0.119034	
total_rooms	0.927253	0.857126	0.918484	0.198050	
total_bedrooms	1.000000	0.873910	0.974725	-0.007682	
population	0.873910	1.000000	0.907222	0.004834	
households	0.974725	0.907222	1.000000	0.013033	
median_income	-0.007682	0.004834	0.013033	1.000000	
median_house_value	0.049454	-0.024650	0.065843	0.688075	

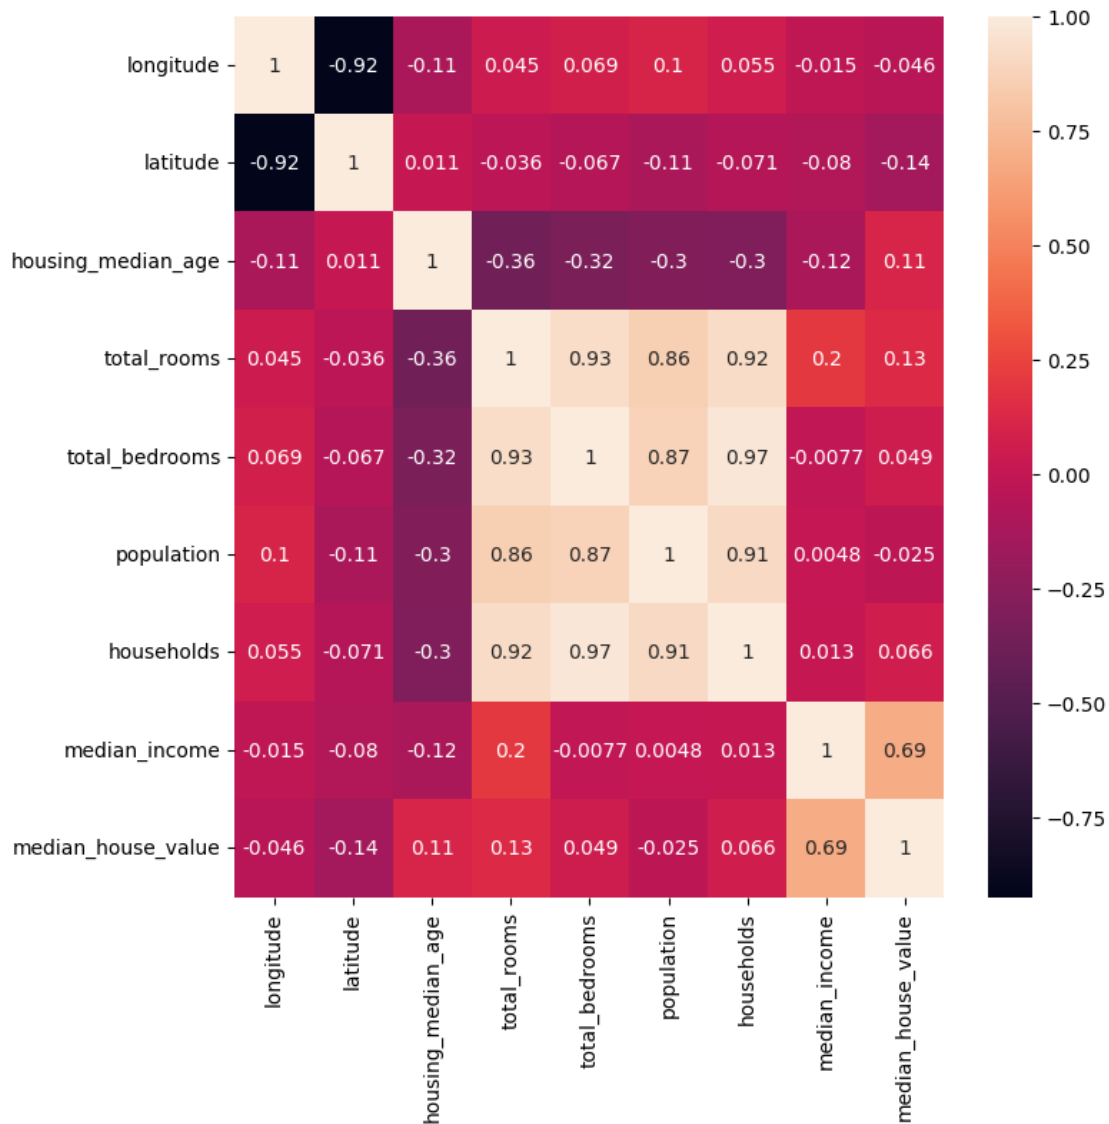
	median_house_value
longitude	-0.045967
latitude	-0.144160
housing_median_age	0.105623
total_rooms	0.134153
total_bedrooms	0.049454
population	-0.024650
households	0.065843
median_income	0.688075
median_house_value	1.000000

<ipython-input-15-3abd71ce2464>:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
corr = housing_df.corr() # data frame correlation function
```

```
[16]: # make the heatmap larger in size
plt.figure(figsize = (8,8))

sns.heatmap(corr, annot=True)
plt.show()
```



```
[17]: # Additionally we noted that several features
      ↪ (total_rooms, total_bedrooms, population, households) have very high
      ↪ correlation to one another,
      # so it's interesting to find out if a removal of a few of them would have any
      ↪ affect on the model performance

      # a new feature that is a ratio of the total rooms to households
      housing_df['rooms_per_household'] = housing_df['total_rooms']/
      ↪ housing_df['households']

      # a new feature that is a ratio of the total bedrooms to the total rooms
```



```

housing_df['bedrooms_per_room'] = housing_df['total_bedrooms']/
↳housing_df['total_rooms']

# a new feature that is a ratio of the population to the households
housing_df['population_per_household']= housing_df['population']/
↳housing_df['households']

# let's combine the latitude and longitude into 1
housing_df['coords'] = housing_df['longitude']/housing_df['latitude']

housing_df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   longitude                             20640 non-null  float64
1   latitude                             20640 non-null  float64
2   housing_median_age                    20640 non-null  float64
3   total_rooms                           20640 non-null  float64
4   total_bedrooms                        20640 non-null  float64
5   population                            20640 non-null  float64
6   households                            20640 non-null  float64
7   median_income                         20640 non-null  float64
8   median_house_value                    20640 non-null  float64
9   ocean_proximity                       20640 non-null  object
10  rooms_per_household                   20640 non-null  float64
11  bedrooms_per_room                     20640 non-null  float64
12  population_per_household              20640 non-null  float64
13  coords                               20640 non-null  float64
dtypes: float64(13), object(1)
memory usage: 2.2+ MB

```

```

[18]: # remove total_rooms, households, total bedrooms, popluation, longitude,
↳latitude
housing_df = housing_df.drop('total_rooms', axis=1)
housing_df = housing_df.drop('households', axis=1)
housing_df = housing_df.drop('total_bedrooms', axis=1)
housing_df = housing_df.drop('population', axis=1)
housing_df = housing_df.drop('longitude', axis=1)
housing_df = housing_df.drop('latitude', axis=1)

housing_df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639

```

Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
0	housing_median_age	20640 non-null	float64
1	median_income	20640 non-null	float64
2	median_house_value	20640 non-null	float64
3	ocean_proximity	20640 non-null	object
4	rooms_per_household	20640 non-null	float64
5	bedrooms_per_room	20640 non-null	float64
6	population_per_household	20640 non-null	float64
7	coords	20640 non-null	float64

dtypes: float64(7), object(1)

memory usage: 1.3+ MB

[19]: *#Heatmap after removing correlation*


```
corr = housing_df.corr()

#make the heatmap larger in size
plt.figure(figsize = (7,7))

sns.heatmap(corr, annot=True)
plt.show()
```

<ipython-input-19-1264607259b1>:3: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
corr = housing_df.corr()
```



```
[20]: #Encoding categorical data
# Most ML algorithms can only learn from numeric data (it's all Math) so
↳ categorical data must be encoded (i.e. converted) to numeric data

# Let's review our data types again; showing that ocean_proximity is the only
↳ categorical data
housing_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
#   ...
```

```

---  -----  -----  -----
0   housing_median_age      20640 non-null float64
1   median_income           20640 non-null float64
2   median_house_value      20640 non-null float64
3   ocean_proximity         20640 non-null object
4   rooms_per_household     20640 non-null float64
5   bedrooms_per_room       20640 non-null float64
6   population_per_household 20640 non-null float64
7   coords                  20640 non-null float64
dtypes: float64(7), object(1)
memory usage: 1.3+ MB

```

```
[21]: # let's see the unique categories for OCEAN_PROXIMITY
housing_df.ocean_proximity.unique()
```

```
[21]: array(['NEAR BAY', '<1H OCEAN', 'INLAND', 'NEAR OCEAN', 'ISLAND'],
      dtype=object)
```

```
[22]: # let's count
housing_df["ocean_proximity"].value_counts()
```

```
[22]: <1H OCEAN      9136
      INLAND      6551
      NEAR OCEAN   2658
      NEAR BAY    2290
      ISLAND        5
      Name: ocean_proximity, dtype: int64
```

```
[23]: # Let's see how the Panda's get_dummies() function works (generates new columns
      ↳ based on the possible options)
print(pd.get_dummies(housing_df['ocean_proximity']))
```

	<1H OCEAN	INLAND	ISLAND	NEAR BAY	NEAR OCEAN
0	0	0	0	1	0
1	0	0	0	1	0
2	0	0	0	1	0
3	0	0	0	1	0
4	0	0	0	1	0
...
20635	0	1	0	0	0
20636	0	1	0	0	0
20637	0	1	0	0	0
20638	0	1	0	0	0
20639	0	1	0	0	0

[20640 rows x 5 columns]

```
[24]: # let's replace the OCEAN_PROXIMITY column using get_dummies()
housing_df_encoded = pd.get_dummies(data=housing_df,
↳ columns=['ocean_proximity'])

# print the first few observations; notice the old OCEAN_PROXIMITY column is
↳ gone
housing_df_encoded.head()
```

```
[24]: housing_median_age  median_income  median_house_value  rooms_per_household \
0                41.0           8.3252           452600.0           6.984127
1                21.0           8.3014           358500.0           6.238137
2                52.0           7.2574           352100.0           8.288136
3                52.0           5.6431           341300.0           5.817352
4                52.0           3.8462           342200.0           6.281853

bedrooms_per_room  population_per_household  coords \
0          0.146591                2.555556 -3.226769
1          0.155797                2.109842 -3.228209
2          0.129516                2.802260 -3.229590
3          0.184458                2.547945 -3.229855
4          0.172096                2.181467 -3.229855

ocean_proximity_<1H OCEAN  ocean_proximity_INLAND  ocean_proximity_ISLAND \
0                0                0                0
1                0                0                0
2                0                0                0
3                0                0                0
4                0                0                0

ocean_proximity_NEAR BAY  ocean_proximity_NEAR OCEAN
0                1                0
1                1                0
2                1                0
3                1                0
4                1                0
```

```
[25]: #Train the model
import sklearn
from sklearn.model_selection import train_test_split

# remove spaces from column names and convert all to lowercase and remove
↳ special characters as it could cause issues in the future
housing_df_encoded.columns = [c.lower().replace(' ', '_').replace('<', '_') for
↳ c in housing_df_encoded.columns]

# Split target variable and feature variables
```

```

X = housing_df_encoded[['housing_median_age',
↪ 'median_income', 'bedrooms_per_room', 'population_per_household', 'coords', 'ocean_proximity__1.
↪ 'ocean_proximity_inland', 'ocean_proximity_island', 'ocean_proximity_near_bay', 'ocean_proximi
y = housing_df_encoded['median_house_value']

print(X)

```

	housing_median_age	median_income	bedrooms_per_room	\
0	41.0	8.3252	0.146591	
1	21.0	8.3014	0.155797	
2	52.0	7.2574	0.129516	
3	52.0	5.6431	0.184458	
4	52.0	3.8462	0.172096	
...	
20635	25.0	1.5603	0.224625	
20636	18.0	2.5568	0.215208	
20637	17.0	1.7000	0.215173	
20638	18.0	1.8672	0.219892	
20639	16.0	2.3886	0.221185	

	population_per_household	coords	ocean_proximity__1h_ocean	\
0	2.555556	-3.226769	0	
1	2.109842	-3.228209	0	
2	2.802260	-3.229590	0	
3	2.547945	-3.229855	0	
4	2.181467	-3.229855	0	
...	
20635	2.560606	-3.067123	0	
20636	3.122807	-3.069385	0	
20637	2.325635	-3.074309	0	
20638	2.123209	-3.076845	0	
20639	2.616981	-3.079502	0	

	ocean_proximity_inland	ocean_proximity_island	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	
...	
20635	1	0	
20636	1	0	
20637	1	0	
20638	1	0	
20639	1	0	

	ocean_proximity_near_bay	ocean_proximity_near_ocean
0	1	0
1	1	0
2	1	0
3	1	0
4	1	0
...
20635	0	0
20636	0	0
20637	0	0
20638	0	0
20639	0	0

[20640 rows x 10 columns]

```
[26]: # Split training & test data
# Splitting the data into training and testing sets in numpy arrays
# We train the model with 70% of the samples and test with the remaining 30%
# X -> array with the inputs; y -> array of the outputs
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42,
    ↪shuffle=True, test_size=0.3)

# Confirm how the data was split
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

(14448, 10)

(6192, 10)

(14448,)

(6192,)

```
[27]: #Linear Regression - Model Training
# Use scikit-learn's LinearRegression to train the model on both the training
    ↪and evaluate it on the test sets
from sklearn.linear_model import LinearRegression

# Create a Linear regressor using all the feature variables
reg_model = LinearRegression()

# Train the model using the training sets
reg_model.fit(X_train, y_train)
```

[27]: LinearRegression()

```
[28]: #run the predictions on the training and testing data
y_pred_test = reg_model.predict(X_test)
```

```
[29]: #compare the actual values (ie, target) with the values predicted by the model
pred_test_df = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred_test})

pred_test_df
```

```
[29]:
```

	Actual	Predicted
20046	47700.0	103743.050896
3024	45800.0	92451.250932
15663	500001.0	219490.963844
20484	218600.0	283292.425471
9814	278000.0	244228.861575
...
17505	237500.0	210121.340663
13512	67300.0	74907.098235
10842	218400.0	216609.962950
16559	119400.0	127975.072923
5786	209800.0	202803.254310

[6192 rows x 2 columns]

```
[30]: # Determine accuracy using R^2
# R^2 : R squared is another way to evaluate the performance of a regression
# model.
# 1, means that the model is perfect and 0 means the the model will perform
# poorly.
r2_reg_model_test = round(reg_model.score(X_test, y_test),2)

print("R^2 Test: {}".format(r2_reg_model_test))
```

R^2 Test: 0.56

```
[31]: # try another machine learning algorithm : Random Forest
# Use scikit-learn's Random Forest to train the model on both the training and
# evaluate it on the test sets
from sklearn.ensemble import RandomForestRegressor

# Create a regressor using all the feature variables
rf_model = RandomForestRegressor(n_estimators=10,random_state=10)

# Train the model using the training sets
rf_model.fit(X_train, y_train)
```

```
[31]: RandomForestRegressor(n_estimators=10, random_state=10)
```



```
[32]: #run the predictions on the training and testing data
y_rf_pred_test = rf_model.predict(X_test)
```

```
[33]: #compare the actual values (ie, target) with the values predicted by the model
rf_pred_test_df = pd.DataFrame({'Actual': y_test, 'Predicted': y_rf_pred_test})

rf_pred_test_df
```

```
[33]:
```

	Actual	Predicted
20046	47700.0	47840.0
3024	45800.0	92680.0
15663	500001.0	446000.5
20484	218600.0	265320.0
9814	278000.0	240800.0
...
17505	237500.0	231680.1
13512	67300.0	69680.0
10842	218400.0	203930.0
16559	119400.0	126170.0
5786	209800.0	198160.0

[6192 rows x 2 columns]

```
[34]: # Determine accuracy using R^2
from sklearn.metrics import r2_score, mean_squared_error

score = r2_score(y_test, y_rf_pred_test)


print("R^2 - {}%".format(round(score, 2) * 100))
```

R^2 - 75.0%

```
[35]: # Determine RMSE - Root Mean Squared Error on the test data
print('RMSE on test data: ', mean_squared_error(y_test, y_rf_pred_test)**(0.5))
```

RMSE on test data: 57289.11495447338

```
[36]: # Determine feature importance - random forest algorithm is that it gives you
↳ the 'feature importance' for all the variables in the data
# plot the 6 most important features
plt.figure(figsize=(10,6))
feat_importances = pd.Series(rf_model.feature_importances_, index = X_train.
↳ columns)
feat_importances.nlargest(6).plot(kind='barh');
```



```
[37]: # training data with 5 most important features
train_x_if = X_train[['bedrooms_per_room', 'housing_median_age', 'coords',
    ↪ 'ocean_proximity_inland', 'population_per_household', 'median_income']]
test_x_if = X_test[['bedrooms_per_room', 'housing_median_age', 'coords',
    ↪ 'ocean_proximity_inland', 'population_per_household', 'median_income']]

# create an object of the RandomForestRegressor Model
rf_model_if = RandomForestRegressor(n_estimators=10, random_state=10)

# fit the model with the training data
rf_model_if.fit(train_x_if, y_train)

# predict the target on the test data
predict_test_with_if = rf_model_if.predict(test_x_if)
```

```
[38]: # Root Mean Squared Error on the train and test data
print('RMSE on test data: ', mean_squared_error(y_test,
    ↪ predict_test_with_if)**(0.5))
```

RMSE on test data: 57366.910692045196

```
[39]: pip install xgboost
```

Requirement already satisfied: xgboost in /usr/local/lib/python3.10/dist-packages (2.0.3)
 Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from xgboost) (1.25.2)
 Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages

(from xgboost) (1.11.4)

```
[40]: # Extreme Gradient Boosting (XGBoost) is an open-source library that provides
      ↪ an efficient and effective implementation of the gradient boosting algorithm.
      # Use the scikit-learn wrapper classes: XGBRegressor and XGBClassifier.
```

```
# try another machine learning algorithm : XGBoost
from xgboost import XGBRegressor

xgb_model = XGBRegressor()
```

```
[41]: # Train the model using the training sets
xgb_model.fit(X_train, y_train)
```

```
[41]: XGBRegressor(base_score=None, booster=None, callbacks=None,
                  colsample_bylevel=None, colsample_bynode=None,
                  colsample_bytreet=None, device=None, early_stopping_rounds=None,
                  enable_categorical=False, eval_metric=None, feature_types=None,
                  gamma=None, grow_policy=None, importance_type=None,
                  interaction_constraints=None, learning_rate=None, max_bin=None,
                  max_cat_threshold=None, max_cat_to_onehot=None,
                  max_delta_step=None, max_depth=None, max_leaves=None,
                  min_child_weight=None, missing=nan, monotone_constraints=None,
                  multi_strategy=None, n_estimators=None, n_jobs=None,
                  num_parallel_tree=None, random_state=None, ...)
```

```
[42]: #run the predictions on the training and testing data
y_xgb_pred_test = xgb_model.predict(X_test)
```

```
[43]: #compare the actual values (ie, target) with the values predicted by the model
xgb_pred_test_df = pd.DataFrame({'Actual': y_test, 'Predicted':
      ↪ y_xgb_pred_test})

xgb_pred_test_df
```


```
[43]:
```

	Actual	Predicted
20046	47700.0	66404.914062
3024	45800.0	86681.765625
15663	500001.0	449666.093750
20484	218600.0	262887.281250
9814	278000.0	218322.796875
...
17505	237500.0	227466.500000
13512	67300.0	64712.433594
10842	218400.0	218226.109375
16559	119400.0	123181.968750
5786	209800.0	227016.828125

[6192 rows x 2 columns]

```
[44]: fig= plt.figure(figsize=(8,8))
xgb_pred_test_df = xgb_pred_test_df.reset_index()
xgb_pred_test_df = xgb_pred_test_df.drop(['index'],axis=1)
plt.plot(xgb_pred_test_df[:50])
plt.legend(['Actual value', 'Predicted value'])
```

[44]: <matplotlib.legend.Legend at 0x78b3029be8c0>



```
[45]: from sklearn.metrics import r2_score

score = r2_score(y_test, y_xgb_pred_test)
```

```
print("R^2 - {}".format(round(score, 2) *100))
```

R^2 - 78.0%

```
[46]: # Determine mean square error and root mean square error
from sklearn.metrics import mean_squared_error
import math

mse = mean_squared_error(y_test, y_xgb_pred_test)
rmse = math.sqrt(mean_squared_error(y_test, y_xgb_pred_test))

print(mse)
print(rmse)
```

2939759040.9080276
54219.5448238735

```
[47]: # Calculate mean absolute error(any large error)
from sklearn.metrics import mean_absolute_error

print(mean_absolute_error(y_test, y_xgb_pred_test))
```

36285.050324826894

```
[48]: # We can build and score a model on multiple folds using cross-validation
from sklearn.model_selection import RepeatedKFold
from sklearn.model_selection import cross_val_score

# define model evaluation method
cv = RepeatedKFold(n_splits=10, n_repeats=3, random_state=1)

scores = cross_val_score(xgb_model, X, y, scoring='r2', error_score='raise',
    ↪cv=cv, n_jobs=-1, verbose=1)

#average of all the r2 scores across runs
print(scores.mean())
```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.

0.7850403811484551

[Parallel(n_jobs=-1)]: Done 30 out of 30 | elapsed: 7.0s finished

```
[49]: # determine hyperparameter available for tuning
xgb_model.get_params()
```

```
[49]: {'objective': 'reg:squarederror',
      'base_score': None,
      'booster': None,
      'callbacks': None,
      'colsample_bylevel': None,
      'colsample_bynode': None,
      'colsample_bytree': None,
      'device': None,
      'early_stopping_rounds': None,
      'enable_categorical': False,
      'eval_metric': None,
      'feature_types': None,
      'gamma': None,
      'grow_policy': None,
      'importance_type': None,
      'interaction_constraints': None,
      'learning_rate': None,
      'max_bin': None,
      'max_cat_threshold': None,
      'max_cat_to_onehot': None,
      'max_delta_step': None,
      'max_depth': None,
      'max_leaves': None,
      'min_child_weight': None,
      'missing': nan,
      'monotone_constraints': None,
      'multi_strategy': None,
      'n_estimators': None,
      'n_jobs': None,
      'num_parallel_tree': None,
      'random_state': None,
      'reg_alpha': None,
      'reg_lambda': None,
      'sampling_method': None,
      'scale_pos_weight': None,
      'subsample': None,
      'tree_method': None,
      'validate_parameters': None,
      'verbosity': None}
```

```
[50]: xgb_model_2 = XGBRegressor(
      gamma=0.05,
      learning_rate=0.01,
      max_depth=6,
      n_estimators=1000,
      n_jobs=16,
      objective='reg:squarederror',
```

```

        subsample=0.8,
        scale_pos_weight=0,
        reg_alpha=0,
        reg_lambda=1,
        verbosity=1)

xgb_model_2.fit(X_train, y_train)

#run the predictions on the training and testing data
y_xgb_2_pred_test = xgb_model_2.predict(X_test)

```

```

[51]: # compare the actual values (ie, target) with the values predicted by the model
xgb_2_pred_test_df = pd.DataFrame({'Actual': y_test, 'Predicted':
    ↪ y_xgb_2_pred_test})

xgb_2_pred_test_df

```

```

[51]:
      Actual    Predicted
20046  47700.0  57542.468750
3024   45800.0  90140.296875
15663 500001.0  441852.906250
20484  218600.0  254412.796875
9814   278000.0  240307.781250
...
17505  237500.0  234835.000000
13512   67300.0   64357.855469
10842  218400.0  220460.828125
16559  119400.0  125676.593750
5786   209800.0  208793.187500

```

[6192 rows x 2 columns]

```


[52]: fig= plt.figure(figsize=(8,8))
xgb_2_pred_test_df = xgb_2_pred_test_df.reset_index()
xgb_2_pred_test_df = xgb_2_pred_test_df.drop(['index'],axis=1)
plt.plot(xgb_2_pred_test_df[:50])
plt.legend(['Actual value', 'Predicted value'])

```

```

[52]: <matplotlib.legend.Legend at 0x78b3029bf0d0>

```



```
[53]: from sklearn.metrics import mean_squared_error

mse = np.sqrt(mean_squared_error(y_test, y_xgb_2_pred_test))
print("RMSE: %.2f" % (mse*(1/2.0)))
```

RMSE: 230.63

```
[54]: # Determine accuracy using r^2
r2_xgb_model_2_test = round(xgb_model_2.score(X_test, y_test),2)

print("R^2 Test: {}".format(r2_xgb_model_2_test))
```

R^2 Test: 0.78